Chronos v3

Cluster-aware lazy expert runtime

a EXPERT PAGE LAYOUT · alternative routes

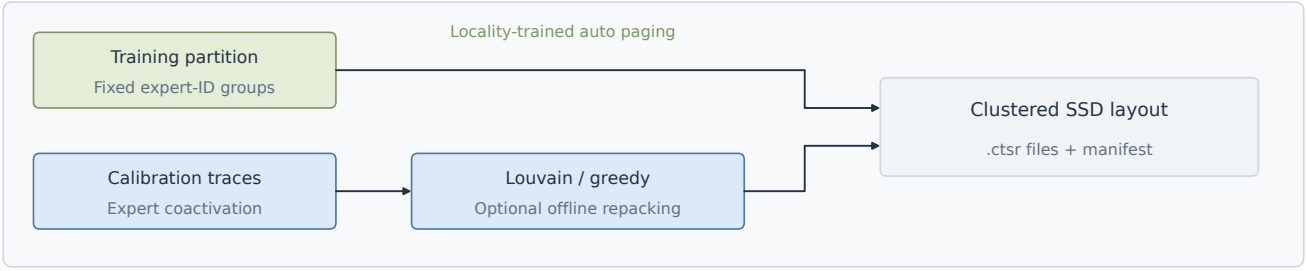

b ASYNC PREDICTION AND WEIGHT DATA PATH · Torch / CUDA

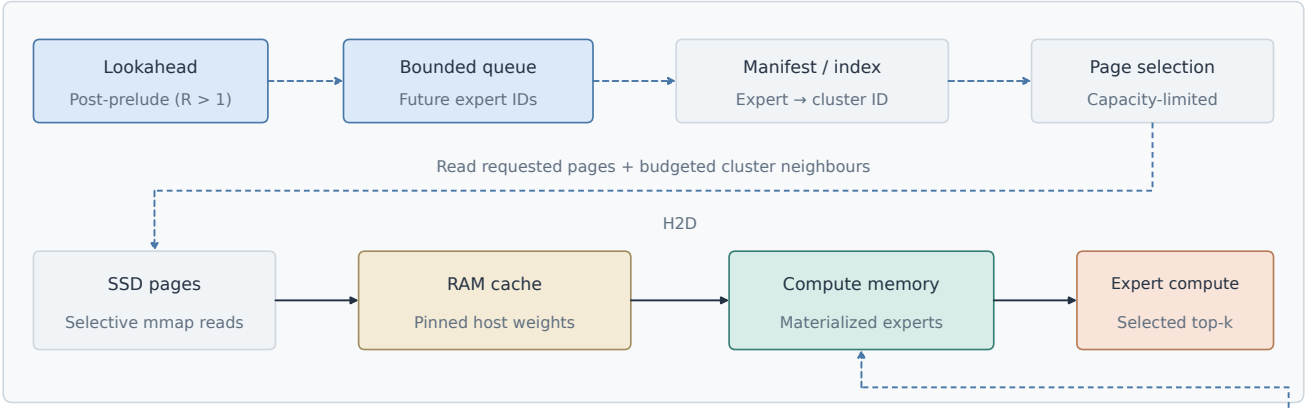

c COLD-MISS POLICY · conditional, not an unconditional replacement

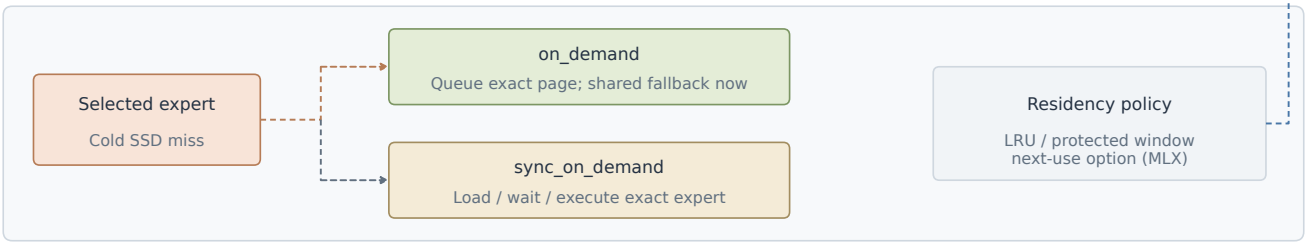

Solid black: expert weight movement. Blue dashes: requests / control. SSD misses are not claimed to be free.

MLX: whole-cluster load / evaluation can precede bounded admission; unified memory, no CUDA H2D.

Locality and packed layout can reduce page demand; improvement requires measurement on the target workload.